

# Project 1 - Classification

## Machine Learning Spring 2021

### Default of Credit card

### Source of Dataset

<https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>  
(<https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>)

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import warnings
6 warnings.filterwarnings("ignore")
7 import seaborn as sns
8 from sklearn.neighbors import KNeighborsClassifier
9 from sklearn.linear_model import LogisticRegression
10 from sklearn.svm import SVC
11 from sklearn.model_selection import train_test_split
12 from sklearn.tree import DecisionTreeClassifier
13 from sklearn import model_selection
14 from sklearn.preprocessing import MinMaxScaler, StandardScaler
15 from sklearn import svm
16 from sklearn.metrics import classification_report
17 from sklearn.metrics import confusion_matrix
18 from sklearn.metrics import accuracy_score
19 from sklearn.model_selection import cross_val_score
20 from sklearn.model_selection import GridSearchCV
21 from sklearn.metrics import recall_score, precision_score, f1_score
22 from sklearn.metrics import precision_recall_curve
23 from sklearn import svm
```

In [2]:

```
1 dcc = pd.read_excel('default of credit card clients.xls', skiprows=1)
2 dcc.drop(['ID'], axis=1, inplace=True)
```

In [3]:

```
1 dcc
```

Out[3]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	...	BILL_AM
0	20000	2	2	1	24	2	2	-1	-1	-2	...	
1	120000	2	2	2	26	-1	2	0	0	0	...	32
2	90000	2	2	2	34	0	0	0	0	0	...	143
3	50000	2	2	1	37	0	0	0	0	0	...	283
4	50000	1	2	1	57	-1	0	-1	0	0	...	209
...	...	...	...	...	...	...	...	...	...	...	...	
29995	220000	1	3	1	39	0	0	0	0	0	...	880
29996	150000	1	3	2	43	-1	-1	-1	-1	0	...	89

In [4]:

```
1 dcc.head()
```

Out[4]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	...
0	20000	2	2	1	24	2	2	-1	-1	-2	...
1	120000	2	2	2	26	-1	2	0	0	0	...
2	90000	2	2	2	34	0	0	0	0	0	...
3	50000	2	2	1	37	0	0	0	0	0	...
4	50000	1	2	1	57	-1	0	-1	0	0	...

5 rows × 24 columns

In [5]:

```
1 dcc.describe()
```

Out[5]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	167484.322667	1.603733	1.853133	1.551867	35.485500	-0.016700
std	129747.661567	0.489129	0.790349	0.521970	9.217904	1.123802
min	10000.000000	1.000000	0.000000	0.000000	21.000000	-2.000000
25%	50000.000000	1.000000	1.000000	1.000000	28.000000	-1.000000
50%	140000.000000	2.000000	2.000000	2.000000	34.000000	0.000000
75%	240000.000000	2.000000	2.000000	2.000000	41.000000	0.000000
max	1000000.000000	2.000000	6.000000	3.000000	79.000000	8.000000

8 rows × 24 columns

In [6]:

```
1 dcc.isnull().sum()
```

Out[6]:

LIMIT_BAL	0
SEX	0
EDUCATION	0
MARRIAGE	0
AGE	0
PAY_0	0
PAY_2	0
PAY_3	0
PAY_4	0
PAY_5	0
PAY_6	0
BILL_AMT1	0
BILL_AMT2	0
BILL_AMT3	0
BILL_AMT4	0
BILL_AMT5	0
BILL_AMT6	0
PAY_AMT1	0

Delete 5.5% of random values, Categorical variables and Age are left alone

In [7]:

```

1 np.random.seed(seed=0)
2 masking_array= np.random.randint(100,size=(dcc.shape[0], 19)) < 94.5
3 masking_array
4

```

Out[7]:

```

array([[ True,  True,  True, ...,  True,  True,  True],
       [ True,  True,  True, ...,  True,  True,  True],
       [ True,  True,  True, ...,  True,  True,  True],
       ...,
       [ True,  True,  True, ...,  True,  True,  True],
       [ True,  True,  True, ...,  True,  True,  True],
       [ True,  True,  True, ...,  True, False,  True]])

```

In [8]:

```

1 dcc[dcc.columns[5:25]]=dcc[dcc.columns[5:25]].where(masking_array, np.nan)

```

In [9]:

```

1 dcc.isna().sum()

```

Out[9]:

```

LIMIT_BAL                0
SEX                      0
EDUCATION                0
MARRIAGE                 0
AGE                     0
PAY_0                   1496
PAY_2                   1490
PAY_3                   1550
PAY_4                   1502
PAY_5                   1437
PAY_6                   1474
BILL_AMT1               1459
BILL_AMT2               1473
BILL_AMT3               1514
BILL_AMT4               1432
BILL_AMT5               1462
BILL_AMT6               1530
PAY_AMT1                1466
PAY_AMT2                1518
PAY_AMT3                1491
PAY_AMT4                1440
PAY_AMT5                1475
PAY_AMT6                1612
default payment next month  1518
dtype: int64

```

**Function to find missing data in percentage**

In [10]:

```

1 def missing_data_percentage(df):
2     x = ['column_name', 'missing_values', 'missing_in_percentage']
3     missing_data = pd.DataFrame(columns=x)
4     columns = dcc.columns
5     for col in columns:
6         iscolumn_name = col
7         ismissing_values = dcc[col].isnull().sum()
8         ismissing_in_percentage = (dcc[col].isnull().sum()/dcc[col].shape[0])*100
9
10    missing_data.loc[len(missing_data)] = [iscolumn_name, ismissing_values, ismissing_in_percentage]
11    print(missing_data.round(2))
12
13

```

In [11]:

```
1 missing_data_percentage(dcc)
```

	column_name	missing_values	missing_in_percentage
0	LIMIT_BAL	0	0.00
1	SEX	0	0.00
2	EDUCATION	0	0.00
3	MARRIAGE	0	0.00
4	AGE	0	0.00
5	PAY_0	1496	4.99
6	PAY_2	1490	4.97
7	PAY_3	1550	5.17
8	PAY_4	1502	5.01
9	PAY_5	1437	4.79
10	PAY_6	1474	4.91
11	BILL_AMT1	1459	4.86
12	BILL_AMT2	1473	4.91
13	BILL_AMT3	1514	5.05
14	BILL_AMT4	1432	4.77
15	BILL_AMT5	1462	4.87
16	BILL_AMT6	1530	5.10
17	PAY_AMT1	1466	4.89
18	PAY_AMT2	1518	5.06
19	PAY_AMT3	1491	4.97
20	PAY_AMT4	1440	4.80
21	PAY_AMT5	1475	4.92
22	PAY_AMT6	1612	5.37
23	default payment next month	1518	5.06

## Data Cleaning

In [12]:

```
1 print(dcc.apply(lambda col: col.unique()))
```

```
LIMIT_BAL      [20000, 120000, 90000, 50000, 500000, 100000,
...
SEX            [2,
1]
EDUCATION      [2, 1, 3, 5, 4, 6,
0]
MARRIAGE       [1, 2, 3,
0]
AGE            [24, 26, 34, 37, 57, 29, 23, 28, 35, 51, 41,
3...
PAY_0          [2.0, -1.0, 0.0, nan, 1.0, -2.0, 3.0, 4.0, 8.
0...
PAY_2          [2.0, 0.0, nan, -2.0, -1.0, 3.0, 5.0, 7.0, 4.
0...
PAY_3          [-1.0, 0.0, 2.0, -2.0, nan, 3.0, 4.0, 6.0, 7.
0...
PAY_4          [-1.0, 0.0, -2.0, 2.0, nan, 3.0, 4.0, 5.0, 7.
0...
PAY_5          [-2.0, 0.0, -1.0, 2.0, nan, 3.0, 5.0, 4.0, 7.
0...
PAY_6          [-2.0, 2.0, 0.0, -1.0, nan, 3.0, 4.0, 6.0, 7.
0...
BILL_AMT1      [3913.0, 2682.0, 29239.0, 46990.0, 8617.0, 64
4...
BILL_AMT2      [3102.0, 1725.0, 14027.0, 48233.0, 5670.0, 57
0...
BILL_AMT3      [689.0, 2682.0, 13559.0, 49291.0, 35835.0, 57
6...
BILL_AMT4      [0.0, 3272.0, 14331.0, 28314.0, 20940.0, 1939
4...
BILL_AMT5      [0.0, 3455.0, 14948.0, 28959.0, 19146.0, nan,
...
BILL_AMT6      [0.0, 3261.0, 15549.0, 29547.0, 19131.0, 2002
4...
PAY_AMT1       [0.0, 1518.0, 2000.0, 2500.0, 55000.0, 380.0,
...
PAY_AMT2       [689.0, 1000.0, 1500.0, 2019.0, 36681.0, 181
5...
PAY_AMT3       [0.0, nan, 1000.0, 1200.0, 10000.0, 657.0, 38
0...
PAY_AMT4       [0.0, 1000.0, 1100.0, 9000.0, 20239.0, 581.0,
...
PAY_AMT5       [0.0, 1000.0, 1069.0, 689.0, 13750.0, 1687.0,
...
PAY_AMT6       [0.0, 2000.0, 5000.0, 1000.0, nan, 800.0, 137
7...
default payment next month      [1.0, 0.0, n
an]
dtype: object
```

### Content of Data

LIMIT\_BAL: Amount of given credit in NT dollars (includes individual and family/supplementary credit

SEX: Gender (1=male, 2=female)

EDUCATION: (0=?, 1=graduate school, 2=university, 3=high school, 4=others, 5=unknown, 6=unknown)

MARRIAGE: Marital status (0=?, 1=married, 2=single, 3=others)

AGE: Age in years

PAY\_0: Repayment status in September, 2005 (-1=pay duly, 1=payment delay for one month, 2=payment delay for two months, ... 8=payment delay for eight months, 9=payment delay for nine months and above)

PAY\_2: Repayment status in August, 2005

PAY\_3: Repayment status in July, 2005

PAY\_4: Repayment status in June, 2005

PAY\_5: Repayment status in May, 2005

PAY\_6: Repayment status in April, 2005 )

BILL\_AMT1: Amount of bill statement in September, 2005 (NT dollar)

BILL\_AMT2: Amount of bill statement in August, 2005 (NT dollar)

BILL\_AMT3: Amount of bill statement in July, 2005 (NT dollar)

BILL\_AMT4: Amount of bill statement in June, 2005 (NT dollar)

BILL\_AMT5: Amount of bill statement in May, 2005 (NT dollar)

BILL\_AMT6: Amount of bill statement in April, 2005 (NT dollar)

PAY\_AMT1: Amount of previous payment in September, 2005 (NT dollar)

PAY\_AMT2: Amount of previous payment in August, 2005 (NT dollar)

PAY\_AMT3: Amount of previous payment in July, 2005 (NT dollar)

PAY\_AMT4: Amount of previous payment in June, 2005 (NT dollar)

PAY\_AMT5: Amount of previous payment in May, 2005 (NT dollar)

PAY\_AMT6: Amount of previous payment in April, 2005 (NT dollar)

default.payment.next.month: Default payment (1=yes, 0=no)

Education type for levels 0,4,5,6 is unknown. We combine and label them as 0

In [13]:

```
1 dcc.loc[dcc.EDUCATION >= 4, 'EDUCATION'] = 0
```

In [14]:

```
1 dcc['EDUCATION'].unique()
```

Out[14]:

```
array([2, 1, 3, 0], dtype=int64)
```

Marriage status for 0 and 3 is unknown. We combine and label them as 0

In [15]:

```
1 dcc.loc[dcc.MARRIAGE == 3, 'MARRIAGE'] = 0
```

In [16]:

```
1 dcc['MARRIAGE'].unique()  
2
```

Out[16]:

```
array([1, 2, 0], dtype=int64)
```

### Filling up missing values

In [17]:

```
1 dcc['PAY_0'].fillna(dcc['PAY_0'].mode()[0],inplace= True)  
2 dcc['PAY_2'].fillna(dcc['PAY_2'].mode()[0],inplace= True)  
3 dcc['PAY_3'].fillna(dcc['PAY_3'].mode()[0],inplace= True)  
4 dcc['PAY_4'].fillna(dcc['PAY_4'].mode()[0],inplace= True)  
5 dcc['PAY_5'].fillna(dcc['PAY_5'].mode()[0],inplace= True)  
6 dcc['PAY_6'].fillna(dcc['PAY_6'].mode()[0],inplace= True)
```

In [18]:

```
1 dcc['PAY_AMT1'].fillna(dcc['PAY_AMT1'].mean(),inplace= True)  
2 dcc['PAY_AMT2'].fillna(dcc['PAY_AMT2'].mean(),inplace= True)  
3 dcc['PAY_AMT3'].fillna(dcc['PAY_AMT3'].mean(),inplace= True)  
4 dcc['PAY_AMT4'].fillna(dcc['PAY_AMT4'].mean(),inplace= True)  
5 dcc['PAY_AMT5'].fillna(dcc['PAY_AMT5'].mean(),inplace= True)  
6 dcc['PAY_AMT6'].fillna(dcc['PAY_AMT6'].mean(),inplace= True)
```

In [19]:

```
1 dcc['BILL_AMT1'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
2 dcc['BILL_AMT2'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
3 dcc['BILL_AMT3'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
4 dcc['BILL_AMT4'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
5 dcc['BILL_AMT5'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
6 dcc['BILL_AMT6'].fillna(dcc['BILL_AMT1'].mean(),inplace= True)  
7
```

In [20]:

```
1 dcc['default payment next month'].fillna(dcc['default payment next month'].mode()[0],inplace= True)
```

In [21]:

```
1 dcc.shape
```

Out[21]:

```
(30000, 24)
```



In [22]:

```
1 dcc.rename(columns={'default payment next month':'Default_Payment'}, inplace=True)
```

In [23]:

```
1 dcc.isna().sum()
```

Out[23]:

LIMIT_BAL	0
SEX	0
EDUCATION	0
MARRIAGE	0
AGE	0
PAY_0	0
PAY_2	0
PAY_3	0
PAY_4	0
PAY_5	0
PAY_6	0
BILL_AMT1	0
BILL_AMT2	0
BILL_AMT3	0
BILL_AMT4	0
BILL_AMT5	0
BILL_AMT6	0
PAY_AMT1	0


In [24]:

```
1 dcc
```

Out[24]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
0	20000	2	2	1	24	2.0	2.0	-1.0	-1.0	-2.0
1	120000	2	2	2	26	-1.0	2.0	0.0	0.0	0.0
2	90000	2	2	2	34	0.0	0.0	0.0	0.0	0.0
3	50000	2	2	1	37	0.0	0.0	0.0	0.0	0.0
4	50000	1	2	1	57	-1.0	0.0	-1.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...
29995	220000	1	3	1	39	0.0	0.0	0.0	0.0	0.0
29996	150000	1	3	2	43	-1.0	-1.0	-1.0	-1.0	0.0
29997	30000	1	2	2	37	4.0	3.0	2.0	-1.0	0.0
29998	80000	1	3	1	41	1.0	-1.0	0.0	0.0	0.0
29999	50000	1	2	1	46	0.0	0.0	0.0	0.0	0.0

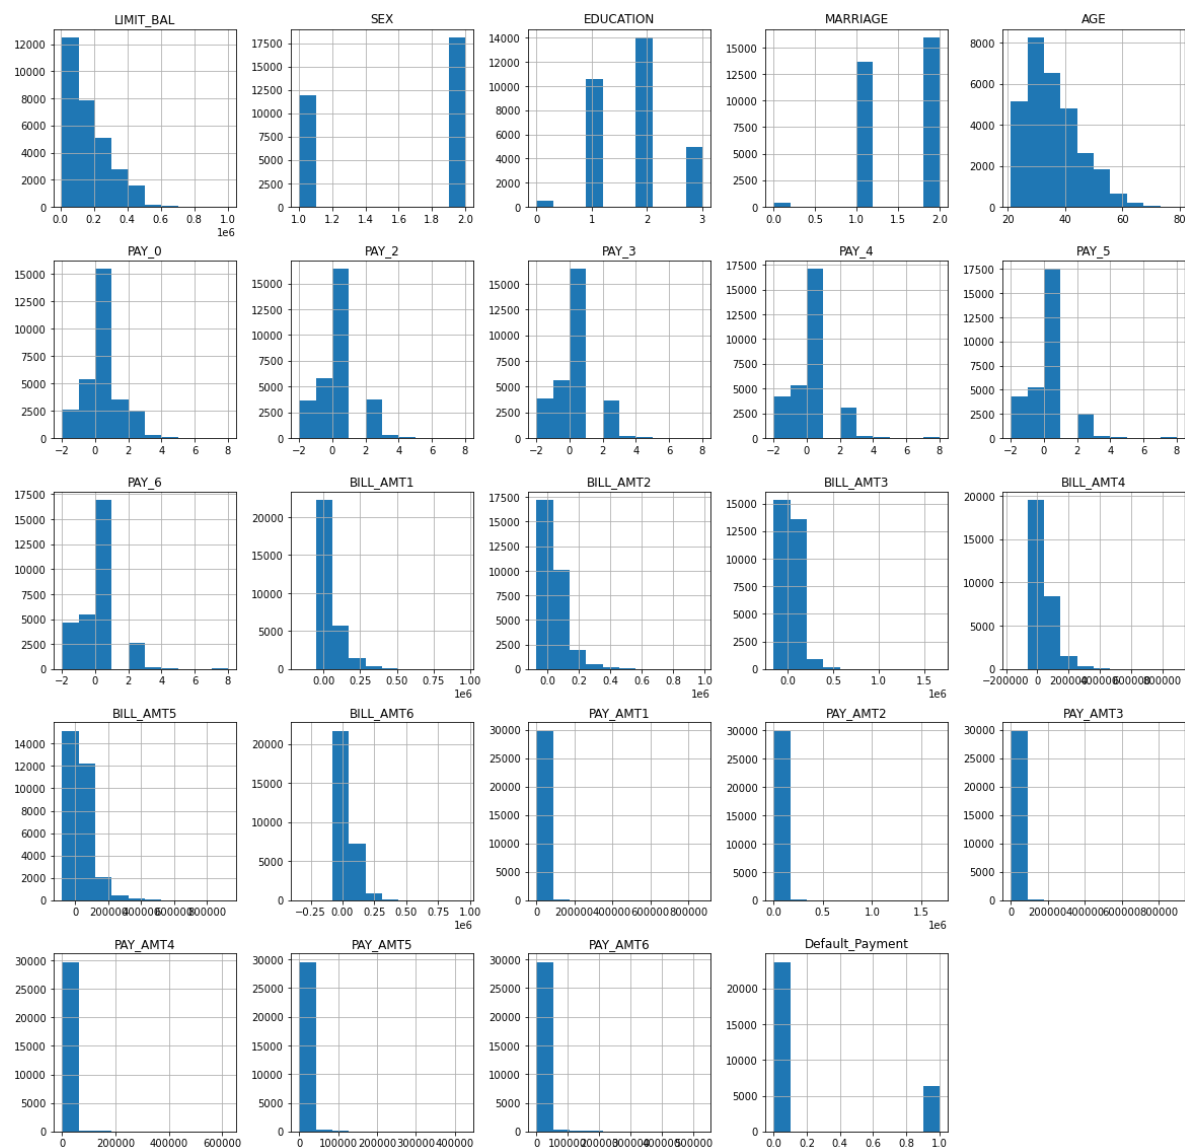
30000 rows × 24 columns



Exploratory Data Analysis


In [25]:

```
1 dcc.hist(figsize=(20,20));
```



In [26]:


```
1 a = (dcc.Default_Payment.value_counts(normalize=True)*100)
2 a.plot.bar(figsize=(9,5))
3 plt.xticks(fontsize=9, rotation=0)
4 plt.yticks(fontsize=9)
5 plt.title("Default_Payment", fontsize=15)
6 for x,y in zip([0,1],a):
7     plt.text(x,y,round(y,2),fontsize=15)
8 plt.show()
```



We can see that the dataset consists of **78.9%** clients are not expected to default payment whereas **21.0%** clients are expected to default the payment. This graph also indicates the our target data column(**Default\_Payment**) is **imbalanced**.

In [27]:

```
1 plt.figure(figsize=(20,20))
2 sns.heatmap(dcc.corr(), annot=True);
```



There is a weak correlation between all the columns of PAY and BILL\_AMT and There is Negative correlation between LIMIT\_BAL and PAY column values

## Data Prepartion for Analysis and classification

### 1. Train Test Split

In [28]:

```
1 X = dcc.drop('Default_Payment',axis =1)
2 y = dcc['Default_Payment']
3 X_train_org, X_test_org, y_train, y_test = train_test_split(X, y, random_state = 0)
```

### 2. Scaling

In [29]:

```
1 scaler = MinMaxScaler()
2 X_train = scaler.fit_transform(X_train_org)
3 X_test = scaler.transform(X_test_org)
```

## Classification Model

### 1. k-nearest neighbors (KNN)

In [30]:


```

1 %matplotlib inline
2 from sklearn.neighbors import KNeighborsClassifier
3
4 knn_training = []
5 knn_testing = []
6 xvalues = range(1,20)
7
8 for i in range(1,20):
9     knn=KNeighborsClassifier(n_neighbors = i)
10    knn.fit(X_train,y_train)
11    train_score_knn=knn.score(X_train,y_train)
12    test_score_knn=knn.score(X_test,y_test)
13    knn_training.append(train_score_knn)
14    knn_testing.append(test_score_knn)
15
16 plt.subplots(figsize = (15,5))
17 plt.plot(xvalues, knn_training, color='g', label = 'Training Score')
18 plt.plot(xvalues, knn_testing, color='r', label = 'Test')
19 plt.xticks(xvalues, range(19))
20 plt.xlabel('Number of neighbor')
21 plt.ylabel('Accuracy')
22 plt.legend()
23
24

```

Out[30]:

&lt;matplotlib.legend.Legend at 0x233135e9400&gt;



In [31]:

```

1 knn=KNeighborsClassifier(n_neighbors = 9)
2 knn.fit(X_train, y_train)
3

```

Out[31]:

KNeighborsClassifier(n\_neighbors=9)

neighbors(k) = 9 is the best parameter for knn model

In [32]:

```
1 print('Training score: {:.3f}'.format(knn.score(X_train, y_train)))
2 print('Testing score: {:.3f}'.format(knn.score(X_test, y_test)))
```

Training score: 0.831

Testing score: 0.807

In [33]:

```
1 X_train.shape
```

Out[33]:

(22500, 23)

### Cross validation scores for KNN Classifier

In [34]:

```
1 knn_CV_scores = cross_val_score(knn, X_train, y_train, cv=7)
2
3 pd.DataFrame({'Train Score - Cross Validation ': knn_CV_scores})
```

Out[34]:

Train Score - Cross Validation	
0	0.805910
1	0.808087
2	0.815184
3	0.809583
4	0.807716
5	0.805227
6	0.808650

### Grid Search on KNN Classifier

In [35]:

```
1 grid_knn_parameters = {'n_neighbors':range(1,20), 'p': [1,2],
2                        'weights': ['uniform','distance'],
3                        'metric': ['euclidean','manhattan']}
4
5 knn_CV = GridSearchCV(KNeighborsClassifier(), grid_knn_parameters, verbose = 1, cv = 7,
6
7 Knn_results = knn_CV.fit(X_train, y_train)
```

Fitting 7 folds for each of 152 candidates, totalling 1064 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 11.2s
[Parallel(n_jobs=-1)]: Done 184 tasks    | elapsed: 55.0s
[Parallel(n_jobs=-1)]: Done 434 tasks    | elapsed: 2.2min
[Parallel(n_jobs=-1)]: Done 784 tasks    | elapsed: 4.4min
[Parallel(n_jobs=-1)]: Done 1064 out of 1064 | elapsed: 6.5min finished
```

In [36]:

```
1 round(knn_CV.score(X_test,y_test),2)
```

Out[36]:

0.81

In [37]:

```
1 print("KNN grid search Best Parameters ")
2 best_parameters_knn=knn_results.best_params_
3 best_parameters_knn
```

KNN grid search Best Parameters

Out[37]:

```
{'metric': 'manhattan', 'n_neighbors': 19, 'p': 1, 'weights': 'distance'}
```

In [38]:

```
1 best_para_Knn = KNeighborsClassifier(metric= 'manhattan', n_neighbors = 19, p = 1, weights = 'distance')
2 best_para_Knn.fit(X_train, y_train)
3 Knn_value_y = best_para_Knn.predict(X_test)
4
5
```

In [39]:

```
1 print('Training score: {:.3f}'.format(best_para_Knn.score(X_train, y_train)))
2 print('Testing score: {:.3f}'.format(best_para_Knn.score(X_test, y_test)))
```

Training score: 1.000

Testing score: 0.810

In [40]:

```
1 print(classification_report(y_pred = Knn_value_y, y_true = y_test))
```

	precision	recall	f1-score	support
0.0	0.83	0.96	0.89	5950
1.0	0.60	0.25	0.35	1550
accuracy			0.81	7500
macro avg	0.71	0.60	0.62	7500
weighted avg	0.78	0.81	0.78	7500

In [41]:

```
1 print(confusion_matrix(y_pred = Knn_value_y, y_true = y_test))
```

```
[[5687 263]
 [1162 388]]
```




In [42]:

```
1 import mglearn
```

In [43]:

```
1 heatmap = mglearn.tools.heatmap(  
2     confusion_matrix(y_pred = Knn_value_y, y_true = y_test), xlabel = 'Predicted label',  
3     ylabel='True label', xticklabels = ['No Default', 'Default'], yticklabels=['No Defau  
4 plt.title("Confusion matrix (KNN)")  
5 plt.gca().invert_yaxis()
```



In [44]:

```
1 Knn_precision_score=precision_score(y_test, best_para_Knn.predict(X_test))  
2 print('Precision score : {:.2f} '.format(Knn_precision_score))
```

Precision score : 0.60

In [45]:

```
1 Knn_recall_score = recall_score(y_test, best_para_Knn.predict(X_test))  
2 print('Recall score : {:.2f} '.format(Knn_recall_score))
```

Recall score : 0.25

In [46]:

```
1 Knn_f1_score = f1_score(y_test,best_para_Knn.predict(X_test))  
2 print('f1 score : {:.2f} '.format(Knn_f1_score))
```

f1 score : 0.35

The precision-recall curve is used for evaluating the performance of binary classification algorithms. It is often used in situations where classes are heavily imbalanced

In [47]:

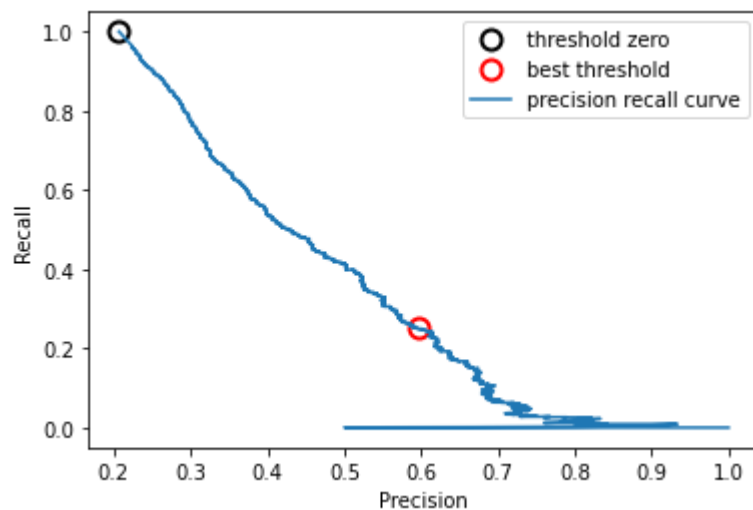
```

1 import mglearn
2
3 %matplotlib inline
4
5 precision, recall, thresholds = precision_recall_curve(y_test, best_para_Knn.predict_proba(X_test))
6
7 close_zero = np.argmin(np.abs(thresholds))
8
9 plt.plot(precision[close_zero], recall[close_zero], 'o', markersize=10, label="threshold zero")
10
11 plt.plot(Knn_precision_score, Knn_recall_score, 'o', markersize=10, label="best threshold")
12
13 plt.plot(precision, recall, label="precision recall curve")
14 plt.xlabel("Precision")
15 plt.ylabel("Recall")
16 plt.legend(loc="best")

```

Out[47]:

&lt;matplotlib.legend.Legend at 0x2330e428ca0&gt;



In [48]:

```

1 Summary_Knn= {'Type': 'K-nearest Neighbors (KNN) Classification Model', 'Training Score': best_para_Knn.score(X_train, y_train)*100,
2               'Testing Score': best_para_Knn.score(X_test, y_test)*100,
3               'f1 Score': f1_score(y_test, best_para_Knn.predict(X_test))};

```

In [49]:

```
1 Summary_Knn
```

Out[49]:

```
{'Type': 'K-nearest Neighbors (KNN) Classification Model',  
'Training Score': 100.0,  
'Testing Score': 81.0,  
'f1 Score': 0.35256701499318494}
```

## 2. Logistic Regression

In [50]:


```
1 c_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]  
2 l1_training = []  
3 l1_testing = []  
4 l2_training = []  
5 l2_testing = []  
6 #As liblinear and saga handle 'l1 and l2' penalty, for small datasets, 'liblinear' is a  
7 for c in c_range:  
8     l1_logistic = LogisticRegression(penalty = 'l1', C = c, solver='liblinear', n_jobs = -1)  
9     l2_logistic = LogisticRegression(penalty = 'l2', C = c, solver = 'lbfgs', n_jobs = -1)  
10    l1_logistic.fit(X_train, y_train)  
11    l2_logistic.fit(X_train, y_train)  
12    l1_training.append(l1_logistic.score(X_train, y_train))  
13    l1_testing.append(l1_logistic.score(X_test, y_test))  
14    l2_training.append(l2_logistic.score(X_train, y_train))  
15    l2_testing.append(l2_logistic.score(X_test, y_test))
```

In [51]:

```

1 plt.plot(c_range, l1_training, label = 'Train score, penalty = l1')
2 plt.plot(c_range, l1_testing, label = 'Test score, penalty = l1')
3 plt.plot(c_range, l2_training, label = 'Train score, penalty = l2')
4 plt.plot(c_range, l2_testing, label = 'Test score, penalty = l2')
5 plt.legend()
6 plt.xlabel('Regularization parameter: C')
7 plt.ylabel('Accuracy')
8 plt.xscale('log')
9

```



According to above graph, C=100 & l2 penalty, test accuracy is best.

In [52]:

```

1 logistic = LogisticRegression(penalty = 'l2', C=100)
2
3 logistic.fit(X_train, y_train)
4
5 print('Training score: {:.3f}'.format(logistic.score(X_train, y_train)))
6 print('Testing score: {:.3f}'.format(logistic.score(X_test, y_test)))

```

Training score: 0.810

Testing score: 0.817

In [53]:

```

1 logistic_CV_scores = cross_val_score(logistic, X_train, y_train, cv=5)
2
3 pd.DataFrame({'Train Score - Cross Validation ': logistic_CV_scores})

```

Out[53]:

Train Score - Cross Validation	
0	0.811333
1	0.812000
2	0.812444
3	0.805111
4	0.811111

In [54]:

```
1 print("Average cross-validation score is : {:.2f}".format(logistic_CV_scores.mean()))
```

Average cross-validation score is : 0.81

## Applying Grid Search with Logistic Regression

In [55]:

```
1 logistic.get_params()
```

Out[55]:

```
{'C': 100,
 'class_weight': None,
 'dual': False,
 'fit_intercept': True,
 'intercept_scaling': 1,
 'l1_ratio': None,
 'max_iter': 100,
 'multi_class': 'auto',
 'n_jobs': None,
 'penalty': 'l2',
 'random_state': None,
 'solver': 'lbfgs',
 'tol': 0.0001,
 'verbose': 0,
 'warm_start': False}
```

In [56]:

```
1 param_grid_logit = { 'max_iter' : range(1,200), 'penalty' : ['l1','l2'],
2                     'C' : [0.001, 0.01, 0.1, 1, 10, 100, 1000]}
3 logit_class_CV = GridSearchCV(estimator = logistic, param_grid = param_grid_logit, cv = 5)
4 GS_results_logit = logit_class_CV.fit(X_train, y_train)
5
6 best_parameters_logit = logit_class_CV.best_params_
7
```

Fitting 5 folds for each of 2786 candidates, totalling 13930 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 56 tasks      | elapsed:    0.3s
[Parallel(n_jobs=-1)]: Done 1200 tasks   | elapsed:    8.0s
[Parallel(n_jobs=-1)]: Done 3200 tasks   | elapsed:   27.2s
[Parallel(n_jobs=-1)]: Done 4768 tasks   | elapsed:   48.7s
[Parallel(n_jobs=-1)]: Done 5218 tasks   | elapsed:   58.6s
[Parallel(n_jobs=-1)]: Done 5768 tasks   | elapsed:  1.2min
[Parallel(n_jobs=-1)]: Done 6874 tasks   | elapsed:  1.6min
[Parallel(n_jobs=-1)]: Done 7624 tasks   | elapsed:  2.2min
[Parallel(n_jobs=-1)]: Done 8930 tasks   | elapsed:  2.8min
[Parallel(n_jobs=-1)]: Done 9880 tasks   | elapsed:  3.7min
[Parallel(n_jobs=-1)]: Done 11386 tasks  | elapsed:  4.5min
[Parallel(n_jobs=-1)]: Done 13104 tasks  | elapsed:  5.7min
[Parallel(n_jobs=-1)]: Done 13930 out of 13930 | elapsed:  6.6min finished
```

In [57]:

```
1 print("Logistic grid search Best score ")
2 GS_results_logit.best_score_
```

Logistic grid search Best score

Out[57]:

0.8110666666666667

In [58]:

```
1 best_parameters = logit_class_CV.best_params_
2 print("Logistic grid search Best parameters: ")
3 best_parameters_logit
```

Logistic grid search Best parameters:

Out[58]:

{ 'C': 1000, 'max\_iter': 90, 'penalty': 'l2' }

**Grid Search on Logistic Regressio with C=1000, max\_iter=90, penalty=l2**

In [59]:

```
1 best_para_logistic = LogisticRegression( C = 1000, max_iter = 90, penalty = 'l2',)
2
3 best_para_logistic.fit(X_train,y_train)
4 logistic_value_y = best_para_logistic.predict(X_test)
5
6 print('Training score: {:.3f}'.format(best_para_logistic.score(X_train, y_train)))
7 print('Testing score: {:.3f}'.format(best_para_logistic.score(X_test, y_test)))
```

Training score: 0.811

Testing score: 0.817

In [60]:

```
1 print(classification_report(y_pred = logistic_value_y, y_true = y_test))
```

	precision	recall	f1-score	support
0.0	0.82	0.98	0.89	5950
1.0	0.71	0.19	0.30	1550
accuracy			0.82	7500
macro avg	0.77	0.58	0.60	7500
weighted avg	0.80	0.82	0.77	7500

In [61]:

```
1 print(confusion_matrix(y_pred =logistic_value_y, y_true = y_test))
```


```
[[5830  120]
 [1256  294]]
```

In [62]:

```

1 heatmap = mglearn.tools.heatmap(
2     confusion_matrix(y_pred = logistic_value_y, y_true = y_test), xlabel = 'Predicted ]
3     ylabel='True label', xticklabels = ['No Default','Default'],yticklabels=['No Defau]
4 plt.title("Confusion matrix (Logistic)")
5 plt.gca().invert_yaxis()

```



In [63]:

```

1 print("Logistic grid search Best Score ")
2 GS_results_logit.best_score_

```

Logistic grid search Best Score

Out[63]:

0.8110666666666667

In [64]:

```

1 logistic_precision_score=precision_score(y_test, best_para_logistic.predict(X_test))
2 print('Precision score : {:.2f} '.format(logistic_precision_score))

```

Precision score : 0.71

In [65]:

```

1 logistic_recall_score = recall_score(y_test, best_para_Knn.predict(X_test))
2 print('Recall score : {:.2f} '.format(logistic_recall_score))

```

Recall score : 0.25

In [66]:

```

1 logistic_f1_score = f1_score(y_test, best_para_logistic.predict(X_test))
2 print('f1 score : {:.2f} '.format(logistic_f1_score))

```

f1 score : 0.30

In [67]:


```

1 precision, recall, thresholds = precision_recall_curve(y_test, best_para_logistic.predict(X_test))
2
3 plt.plot(precision[close_zero], recall[close_zero], 'o', markersize=10,
4          label="threshold zero", fillstyle="none", c='k', mew=2)
5
6 plt.plot(logistic_precision_score, logistic_recall_score, 'o', markersize=10,
7          label="best threshold", fillstyle="none", c='r', mew=2)
8
9 plt.plot(precision, recall, label="Precision recall curve")
10 plt.xlabel("Precision")
11 plt.ylabel("Recall")
12 plt.legend(loc="best")

```

Out[67]:

&lt;matplotlib.legend.Legend at 0x23311bde790&gt;



In [68]:

```

1 Summary_Logistic= {'Type': 'Logistic Regression', 'Train Score': best_para_logistic.score(X_train, y_train)*100,
2                   'Testing Score':best_para_logistic.score(X_test, y_test)*100,
3                   'f1 Score':f1_score(y_test, best_para_logistic.predict(X_test))};

```

In [69]:

```

1 Summary_Logistic

```

Out[69]:

```

{'Type': 'Logistic Regression',
 'Train Score': 81.05333333333333,
 'Testing Score': 81.65333333333334,
 'f1 Score': 0.2993890020366599}

```

### 3. Linear Support Vector Machine Classifier



In [70]:

```

1 from sklearn.svm import LinearSVC,SVC
2
3 c_val = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
4
5 lin_train_score = []
6 lin_test_score = []
7
8 for c in c_val:
9     lin_svm = LinearSVC(C = c)
10    lin_svm.fit(X_train, y_train)
11    lin_train_score.append(lin_svm.score(X_train, y_train))
12    lin_test_score.append(lin_svm.score(X_test, y_test))

```

In [71]:


```

1 plt.plot(c_val, lin_train_score, label = 'Training Score', c = 'g')
2 plt.plot(c_val, lin_test_score, label = 'Testing Score', c = 'r')
3 plt.xscale('log')
4 plt.xlabel('Regularization parameter C')
5 plt.ylabel('Accuracy')
6 plt.legend()

```

Out[71]:

&lt;matplotlib.legend.Legend at 0x23314457790&gt;



Based on above graph C=100, test accuracy is best.

In [72]:

```

1 linear_svm = LinearSVC(C = 100)
2 linear_svm.fit(X_train, y_train)
3
4 print('Training score: {:.3f}'.format(linear_svm.score(X_train, y_train)))
5 print('Testing score: {:.3f}'.format(linear_svm.score(X_test, y_test)))

```

Training score: 0.795

Testing score: 0.802

In [73]:

```

1 linear_svm_cv = cross_val_score(linear_svm, X_train, y_train, cv=5)
2
3 pd.DataFrame({'Train Score - Cross Validation ': linear_svm_cv})

```

Out[73]:

Train Score - Cross Validation	
0	0.789778
1	0.791778
2	0.808667
3	0.791556
4	0.791333

In [74]:

```
1 print("Average cross-validation score: {:.2f}".format(linear_svm_cv.mean()))
```

Average cross-validation score: 0.79

## Applying Grid Search with Linear Support Vector Machine Classifier

In [75]:

```

1 param_linearSVM = { 'max_iter' : range(1,200), 'C' : [ 0.001,0.01, 0.1, 1, 10, 100, 1000]
2
3 CV_linearSVM = GridSearchCV(estimator = lin_svm, param_grid = param_linearSVM ,cv = 5,
4 GS_results_linearSVM = CV_linearSVM.fit(X_train, y_train)
5
6 best_parameters_linearSVM = CV_linearSVM.best_params_
7 print(best_parameters_linearSVM)

```

Fitting 5 folds for each of 1393 candidates, totalling 6965 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 52 tasks      | elapsed: 1.0s
[Parallel(n_jobs=-1)]: Done 352 tasks    | elapsed: 8.0s
[Parallel(n_jobs=-1)]: Done 852 tasks    | elapsed: 20.4s
[Parallel(n_jobs=-1)]: Done 1552 tasks   | elapsed: 35.5s
[Parallel(n_jobs=-1)]: Done 2452 tasks   | elapsed: 1.0min
[Parallel(n_jobs=-1)]: Done 3392 tasks   | elapsed: 1.9min
[Parallel(n_jobs=-1)]: Done 4042 tasks   | elapsed: 3.8min
[Parallel(n_jobs=-1)]: Done 4792 tasks   | elapsed: 5.6min
[Parallel(n_jobs=-1)]: Done 5642 tasks   | elapsed: 7.6min
[Parallel(n_jobs=-1)]: Done 6592 tasks   | elapsed: 10.0min
[Parallel(n_jobs=-1)]: Done 6965 out of 6965 | elapsed: 11.6min finished

```

{'C': 10, 'max\_iter': 128}

In [76]:

```
1 print("Best score : Linear SVM grid search ")
2 GS_results_linearSVM.best_score_
```

Best score : Linear SVM grid search

Out[76]:

0.8086666666666666

In [77]:

```
1 print("Best parameters : Linear SVM grid search ")
2 best_parameters_linearSVM
```

Best parameters : Linear SVM grid search

Out[77]:

{ 'C': 10, 'max\_iter': 128 }

**GridSearch for Linear SVM Classification with C=10 and max\_iter=128**

In [78]:

```
1 best_para_lin_SVM = LinearSVC(C = 10,max_iter = 128)
2 best_para_lin_SVM.fit(X_train, y_train)
3 SVM_value_y = best_para_lin_SVM.predict(X_test)
4
5 print('Training score: {:.3f}'.format(best_para_lin_SVM.score(X_train, y_train)))
6 print('Testing score: {:.3f}'.format(best_para_lin_SVM.score(X_test, y_test)))
7
8
```

Training score: 0.796

Testing score: 0.803

In [79]:

```
1 print(classification_report(y_pred = SVM_value_y, y_true = y_test))
```

	precision	recall	f1-score	support
0.0	0.80	0.99	0.89	5950
1.0	0.76	0.07	0.12	1550
accuracy			0.80	7500
macro avg	0.78	0.53	0.51	7500
weighted avg	0.79	0.80	0.73	7500


In [80]:

```
1 print(confusion_matrix(y_pred = SVM_value_y, y_true = y_test))
```

```
[[5917  33]
 [1446 104]]
```

In [81]:

```
1 heatmap = mglearn.tools.heatmap(  
2     confusion_matrix(y_pred = SVM_value_y, y_true = y_test), xlabel = 'Predicted label',  
3     ylabel='True label', xticklabels = ['No Default','Default'],yticklabels=['No Default','Default'],  
4     plt.title("Confusion matrix Linear SVM")  
5     plt.gca().invert_yaxis()  
6  
7
```



In [82]:

```
1 lin_SVM_precision_score=precision_score(y_test, best_para_lin_SVM.predict(X_test))  
2 print('Precision score : {:.2f} '.format(lin_SVM_precision_score))  
3
```

Precision score : 0.76

In [83]:

```
1 lin_SVM_recall_score = recall_score(y_test, best_para_lin_SVM.predict(X_test))  
2 print('Recall score : {:.2f} '.format(lin_SVM_recall_score))
```

Recall score : 0.07

In [84]:

```
1 lin_SVM_f1_score = f1_score(y_test, best_para_lin_SVM.predict(X_test))  
2 print('f1 score : {:.2f} '.format(lin_SVM_f1_score))
```

f1 score : 0.12

In [85]:

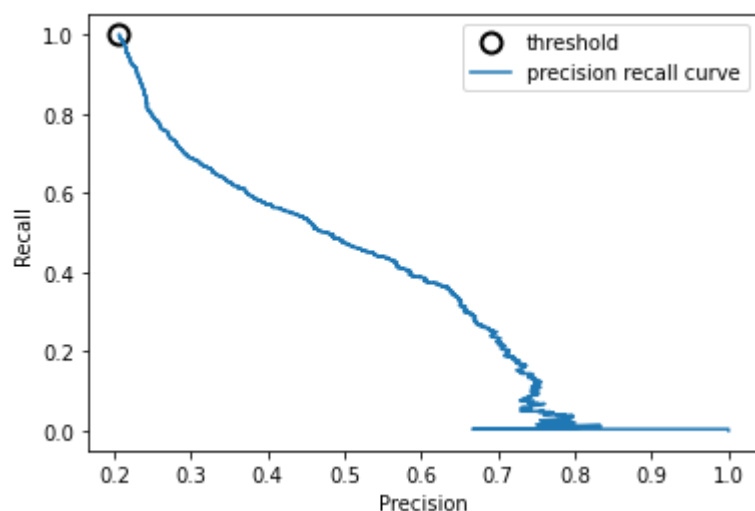
```

1 import mglearn
2 %matplotlib inline
3
4 precision, recall, thresholds = precision_recall_curve(y_test, best_para_lin_SVM.decisi
5
6 plt.plot(precision[close_zero], recall[close_zero], 'o', markersize=10,
7          label="threshold", fillstyle="none", c='k', mew=2)
8
9 plt.plot(precision, recall, label="precision recall curve")
10 plt.xlabel("Precision")
11 plt.ylabel("Recall")
12 plt.legend(loc="best")
13
14
15

```

Out[85]:

&lt;matplotlib.legend.Legend at 0x23311f8afa0&gt;



In [86]:

```

1 Summary_lin_SVM= {'Type': 'Linear SVM', 'Train Score': best_para_lin_SVM.score(X_train,
2                                     'Testing Score':best_para_lin_SVM.score(X_test, y_test)*100,
3                                     'f1 Score':f1_score(y_test, best_para_lin_SVM.predict(X_test))});

```

In [87]:

```
1 Summary_lin_SVM
```

Out[87]:

```
{'Type': 'Linear SVM',
 'Train Score': 79.63555555555556,
 'Testing Score': 80.28,
 'f1 Score': 0.12329579134558386}
```

## 4.Kerenilzed Support Vector Machine (rbf, poly, and linear)

Reducing sample size to 1000 samples with subsampling, map the training samples with random feature mapping to obtain training set and train linear SVMs in parallel to get a unified model on the training set.

In [88]:

```
1 dcc_k = dcc.sample(n = 1000, random_state= 0)
2
3 dcc_k.shape
```

Out[88]:

```
(1000, 24)
```

In [89]:

```
1 dcc_k.head()
```

Out[89]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
<b>8225</b>	20000	1	1	2	33	1.0	2.0	0.0	2.0	2.0
<b>10794</b>	20000	2	2	2	35	0.0	0.0	2.0	0.0	0.0
<b>9163</b>	230000	2	1	1	44	1.0	-1.0	-1.0	-1.0	-1.0
<b>26591</b>	100000	1	2	1	42	0.0	0.0	0.0	0.0	0.0
<b>6631</b>	150000	1	1	2	29	-2.0	-2.0	-2.0	-2.0	-2.0

5 rows × 24 columns

In [90]:

```
1 X_k = dcc_k.drop(['Default_Payment'],axis =1)
2
3 y_k = dcc_k['Default_Payment']
4
5 X_train_org_k, X_test_org_k, y_train_k, y_test_k = train_test_split(X_k, y_k, random_st
```

Scaling the small sample

In [91]:

```

1 from sklearn.preprocessing import MinMaxScaler,StandardScaler
2 scaler_new = MinMaxScaler()
3 X_train_k = scaler_new.fit_transform(X_train_org_k)
4 X_test_k = scaler_new.transform(X_test_org_k)
5

```

In [92]:

```

1 pd.DataFrame(X_train_k).head()

```

Out[92]:

	0	1	2	3	4	5	6	7	8	9	...	
0	0.229730	0.0	0.666667	1.0	0.458333	0.166667	0.125	0.000000	0.333333	0.000000	...	0
1	0.351351	0.0	0.666667	1.0	0.125000	0.333333	0.250	0.285714	0.333333	0.285714	...	0
2	0.081081	0.0	0.333333	1.0	0.166667	0.333333	0.250	0.285714	0.666667	0.285714	...	0
3	0.148649	0.0	0.333333	1.0	0.250000	0.333333	0.250	0.285714	0.333333	0.285714	...	0
4	0.391892	1.0	0.333333	1.0	0.104167	0.333333	0.250	0.285714	0.333333	0.285714	...	0

5 rows × 23 columns

In [93]:

```

1 pd.DataFrame(X_test_k).head()

```

Out[93]:

	0	1	2	3	4	5	6	7	8	9	...	
0	0.148649	1.0	0.333333	0.5	0.458333	0.333333	0.25	0.285714	0.333333	0.285714	...	0.0
1	0.000000	1.0	0.666667	1.0	0.125000	0.500000	0.50	0.571429	0.333333	0.285714	...	0.0
2	0.554054	0.0	0.333333	1.0	0.229167	0.333333	0.00	0.000000	0.000000	0.000000	...	0.0
3	0.054054	1.0	0.666667	0.5	0.395833	0.333333	0.25	0.285714	0.333333	0.285714	...	0.0
4	0.202703	0.0	0.333333	0.5	0.187500	0.166667	0.25	0.285714	0.333333	0.285714	...	0.0

5 rows × 23 columns

In [94]:

```

1 c_range = [0.001,0.01, 0.1, 1, 10,100]
2 k1_train_score = []
3 k1_test_score = []
4 for C in c_range:
5     kernal_new = svm.SVC(kernel = 'linear', C=C)
6     kernal_new.fit(X_train_k,y_train_k)
7     k1_train_score.append(kernal_new.score(X_train_k,y_train_k))
8     k1_test_score.append(kernal_new.score(X_test_k, y_test_k))

```

In [95]:

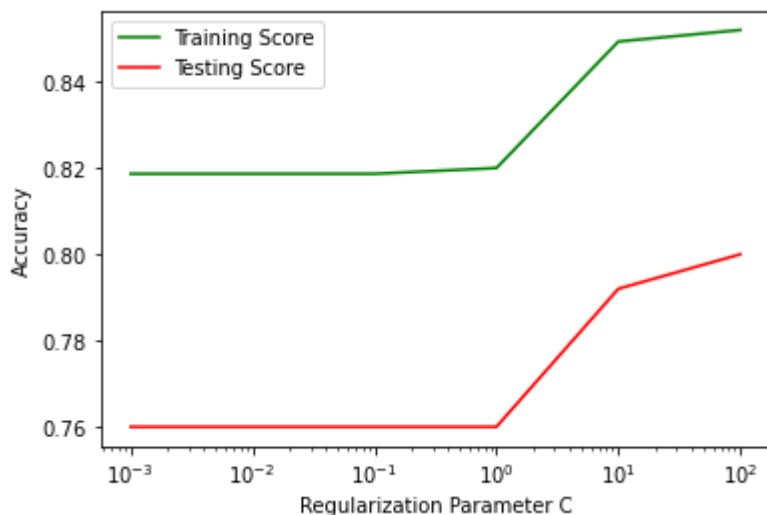
```

1 plt.plot(c_range, k1_train_score, label = 'Training Score', c = 'g')
2 plt.plot(c_range, k1_test_score, label = 'Testing Score', c='r')
3 plt.xscale('log')
4 plt.xlabel('Regularization Parameter C')
5 plt.ylabel('Accuracy')
6
7 plt.legend()

```

Out[95]:

&lt;matplotlib.legend.Legend at 0x23311123ca0&gt;



In [96]:

```

1 from sklearn.svm import LinearSVC,SVC
2 kernal_new = svm.SVC(kernel = 'linear', C=10)
3 kernal_new.fit(X_train_k, y_train_k)
4 print("train score {:.3f}".format(kernal_new.score(X_train_k, y_train_k)))
5 print("test score: {:.3f}".format(kernal_new.score(X_test_k, y_test_k)))
6

```

train score 0.849

test score: 0.792

In [97]:

```

1 from sklearn import svm
2 from sklearn.svm import SVC
3 c_range = [0.001,0.01, 0.1, 1, 10,100]
4 k2_train_score = []
5 k2_test_score = []
6 for C in c_range:
7     kernal_new2 = svm.SVC(kernel = 'poly', C=C)
8     kernal_new2.fit(X_train_k,y_train_k)
9     k2_train_score.append(kernal_new2.score(X_train_k,y_train_k))
10    k2_test_score.append(kernal_new2.score(X_test_k, y_test_k))

```



In [98]:

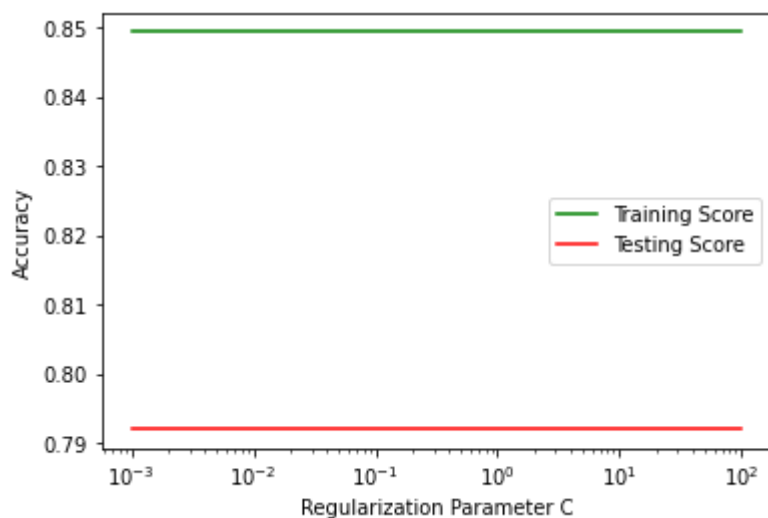
```

1 plt.plot(c_range, k2_train_score, label = 'Training Score', c = 'g')
2 plt.plot(c_range, k2_test_score, label = 'Testing Score', c='r')
3 plt.xscale('log')
4 plt.xlabel('Regularization Parameter C')
5 plt.ylabel('Accuracy')
6
7 plt.legend()

```

Out[98]:

&lt;matplotlib.legend.Legend at 0x23311203d90&gt;



In [99]:

```

1 from sklearn.svm import LinearSVC,SVC
2 kernal_new2 = svm.SVC(kernel = 'poly', C=10)
3 kernal_new2.fit(X_train_k, y_train_k)
4 print("train score {:.3f}".format(kernal_new2.score(X_train_k, y_train_k)))
5 print("test score: {:.3f}".format(kernal_new2.score(X_test_k, y_test_k)))
6

```

train score 0.892

test score: 0.796

In [100]:

```

1 from sklearn import svm
2 from sklearn.svm import SVC
3 c_range = [0.001,0.01, 0.1, 1, 10,100]
4 k3_train_score = []
5 k3_test_score = []
6 for C in c_range:
7     kernal_new3 = svm.SVC(kernel = 'rbf', C=C)
8     kernal_new3.fit(X_train_k,y_train_k)
9     k3_train_score.append(kernal_new3.score(X_train_k,y_train_k))
10    k3_test_score.append(kernal_new3.score(X_test_k, y_test_k))

```

In [101]:

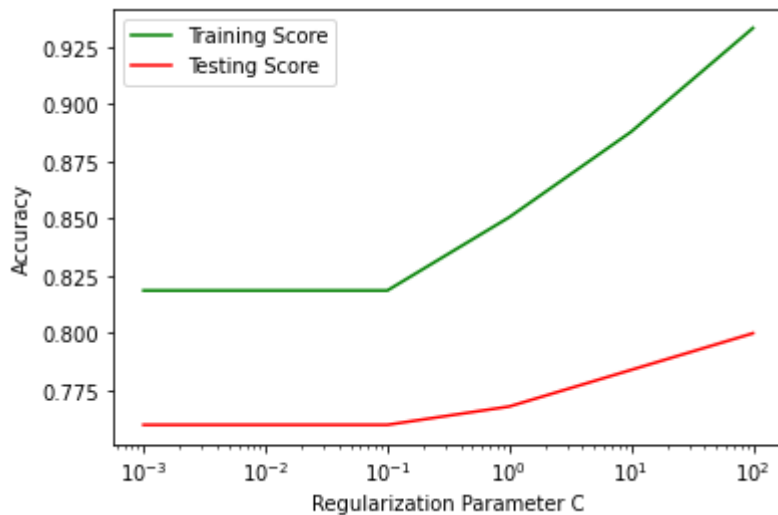
```

1 plt.plot(c_range, k3_train_score, label = 'Training Score', c = 'g')
2 plt.plot(c_range, k3_test_score, label = 'Testing Score', c='r')
3 plt.xscale('log')
4 plt.xlabel('Regularization Parameter C')
5 plt.ylabel('Accuracy')
6
7 plt.legend()

```

Out[101]:

&lt;matplotlib.legend.Legend at 0x23311fa41c0&gt;



In [102]:

```

1 from sklearn.svm import LinearSVC, SVC
2 kernal_new3 = svm.SVC(kernel = 'rbf', C=10)
3 kernal_new3.fit(X_train_k, y_train_k)
4 print("train score {:.3f}".format(kernal_new3.score(X_train_k, y_train_k)))
5 print("test score: {:.3f}".format(kernal_new3.score(X_test_k, y_test_k)))

```

train score 0.888

test score: 0.784

### GridSearch for Kerenilzed Support Vector Machine (rbf, poly, and linear)

In [103]:

```

1 best_score=0
2
3 for gamma in [0.001, 0.01, 0.1, 1, 10, 100]:
4     for C in [0.001, 0.01, 0.1, 1, 10, 100]:
5
6
7         svm = SVC(gamma = gamma, C=C)
8
9         CV_scores = cross_val_score(svm, X_train_k, y_train_k, cv=5)
10
11         CV_score = np.mean(CV_scores)
12
13         if CV_score > best_score:
14             best_score = CV_score
15             best_parameters = {'C': C, 'gamma': gamma}
16
17 # rebuild a model on the combined training and validation set
18 svm = SVC(**best_parameters)
19 svm.fit(X_train_k, y_train_k)

```

Out[103]:

SVC(C=10, gamma=0.1)

In [104]:

```

1 kernelSVC_parameters = {'C':[0.001, 0.01, 0.1, 1, 10, 100], 'gamma':[0.001, 0.01, 0.1, 1, 10, 100]}

```

In [105]:

```

1 print(classification_report(y_pred = SVM_value_y, y_true = y_test))

```

	precision	recall	f1-score	support
0.0	0.80	0.99	0.89	5950
1.0	0.76	0.07	0.12	1550
accuracy			0.80	7500
macro avg	0.78	0.53	0.51	7500
weighted avg	0.79	0.80	0.73	7500

In [106]:

```
1 from sklearn.model_selection import GridSearchCV
2
3 KernelSVC = SVC()
4 GS_KernelSVC = GridSearchCV(KernelSVC, kernelSVC_parameters, cv = 5, return_train_score=
5 GS_KernelSVC.fit(X_train_k,y_train_k)
6
```

Out[106]:

```
GridSearchCV(cv=5, estimator=SVC(), n_jobs=-1,
             param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100],
                         'gamma': [0.001, 0.01, 0.1, 1, 10, 100],
                         'kernel': ['rbf', 'poly', 'linear']},
             return_train_score=True)
```

In [107]:

```
1 print("Best score : KernelSVM grid search ")
2 round(GS_KernelSVC.best_score_,2)
```

Best score : KernelSVM grid search

Out[107]:

0.84

In [108]:

```
1 print("Best parameters- KernelSVM grid search ")
2 GS_KernelSVC.best_params_
```

Best parameters- KernelSVM grid search

Out[108]:

{'C': 0.1, 'gamma': 1, 'kernel': 'poly'}

## kernel = poly

In [109]:

```
1 best_para_svm_poly = SVC(C = 0.1, gamma = 1, kernel = 'poly', verbose = 1)
2
3 best_para_svm_poly.fit(X_train_k,y_train_k)
4 SVM_value_y_poly = best_para_svm_poly.predict(X_test_k)
5 print('Training score: {:.3f}'.format(best_para_svm_poly.score(X_train_k, y_train_k)))
6 print('Testing score: {:.3f}'.format(best_para_svm_poly.score(X_test_k, y_test_k)))
7
8
```

[LibSVM]Training score: 0.856

Testing score: 0.768

In [110]:

```
1 print(classification_report(SVM_value_y_poly,y_test_k))
```

	precision	recall	f1-score	support
0.0	0.98	0.78	0.87	240
1.0	0.10	0.60	0.17	10
accuracy			0.77	250
macro avg	0.54	0.69	0.52	250
weighted avg	0.94	0.77	0.84	250

In [111]:

```
1 import mglearn
```


In [112]:

```
1 print(confusion_matrix(y_pred = SVM_value_y_poly,y_true = y_test_k))
```

```
[[186  4]
 [ 54  6]]
```

In [113]:

```
1 %matplotlib inline
2
3 heatmap = mglearn.tools.heatmap(
4     confusion_matrix(y_pred = SVM_value_y_poly, y_true = y_test_k), xlabel = 'Predicted label',
5     ylabel='True label', xticklabels = ['No Default','Default'], yticklabels=['No Default','Default'],
6     plt.title("Confusion matrix (poly)")
7     plt.gca().invert_yaxis()
```



In [114]:

```

1 Kernel_poly_precision_score=precision_score(y_test_k,best_para_svm_poly.predict(X_test_
2 print('Precision score : {:.2f} '.format(Kernel_poly_precision_score))
3

```

Precision score : 0.60

In [115]:

```

1 Kernel_poly_f1_score=f1_score(y_test_k, best_para_svm_poly.predict(X_test_k))
2 print('f1 Score : {:.2f} '.format(Kernel_poly_f1_score))
3

```

f1 Score : 0.17

In [116]:

```

1 Kernel_poly_recall_score=recall_score(y_test_k, best_para_svm_poly.predict(X_test_k))
2 print('Recall score : {:.2f} '.format(Kernel_poly_recall_score))
3

```

Recall score : 0.10

In [117]:

```

1 Summary_Kernelized_poly= {'Type': 'Kernalized poly ', 'Train Score': best_para_svm_poly
2 'Testing Score':best_para_svm_poly.score(X_test_k, y_test_k)*100,
3 'f1 Score':f1_score(y_test_k, best_para_svm_poly.predict(X_test_k))};

```

In [118]:

```

1 Summary_Kernelized_poly

```

Out[118]:

```

{'Type': 'Kernalized poly ',
 'Train Score': 85.6,
 'Testing Score': 76.8,
 'f1 Score': 0.17142857142857143}

```

## Kernel = rbf

In [119]:

```

1 best_para_svm_rbf = SVC(C = 0.1, gamma = 1, kernel = 'rbf', verbose = 1)
2
3 best_para_svm_rbf.fit(X_train_k,y_train_k)
4 SVM_value_y_rbf = best_para_svm_rbf.predict(X_test_k)
5 print('Training score: {:.3f}'.format(best_para_svm_rbf.score(X_train_k, y_train_k)))
6 print('Testing score: {:.3f}'.format(best_para_svm_rbf.score(X_test_k, y_test_k)))
7

```

[LibSVM]Training score: 0.819  
 Testing score: 0.760

In [120]:

```
1 print(classification_report(SVM_value_y_rbf,y_test_k))
```

	precision	recall	f1-score	support
0.0	1.00	0.76	0.86	250
1.0	0.00	0.00	0.00	0
accuracy			0.76	250
macro avg	0.50	0.38	0.43	250
weighted avg	1.00	0.76	0.86	250

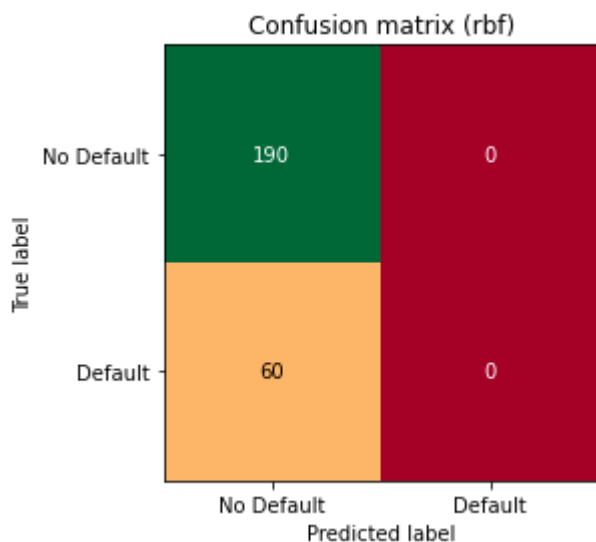
In [121]:

```
1 print(confusion_matrix(y_pred = SVM_value_y_rbf,y_true= y_test_k))
```

```
[[190  0]
 [ 60  0]]
```

In [122]:

```
1 %matplotlib inline
2
3 heatmap = mglearn.tools.heatmap(confusion_matrix(y_pred = SVM_value_y_rbf, y_true = y_t
4 plt.title("Confusion matrix (rbf)")
5 plt.gca().invert_yaxis()
6
```



In [123]:

```
1 Kernel_rbf_precision_score=precision_score(y_test_k,best_para_svm_rbf.predict(X_test_k)
2 print('Precision score : {:.2f} '.format(precision_score(y_test_k,best_para_svm_rbf.pre
3
```

Precision score : 0.00

In [124]:

```

1 Kernel_rbf_f1_score=f1_score(y_test_k, best_para_svm_rbf.predict(X_test_k))
2 print('f1 Score : {:.2f} '.format(Kernel_rbf_f1_score))
3

```

f1 Score : 0.00

In [125]:

```

1 Kernel_rbf_recall_score=recall_score(y_test_k, best_para_svm_rbf.predict(X_test_k))
2 print('Recall score : {:.2f} '.format(Kernel_rbf_recall_score))
3

```

Recall score : 0.00

In [126]:

```

1 Summary_Kernelized_rbf= {'Type': 'Kernalized rbf ', 'Train Score': best_para_svm_rbf.score(X_train_k, y_train_k),
2                           'Testing Score':best_para_svm_rbf.score(X_test_k, y_test_k)*100,
3                           'f1 Score':f1_score(y_test_k, best_para_svm_rbf.predict(X_test_k))};

```

In [127]:

```

1 Summary_Kernelized_rbf

```

Out[127]:

```

{'Type': 'Kernalized rbf ',
 'Train Score': 81.86666666666666,
 'Testing Score': 76.0,
 'f1 Score': 0.0}

```

## Kernel = linear

In [128]:

```

1 best_para_svm_linear = SVC(C = 0.1, cache_size = 200, gamma = 1, kernel = 'linear', ver
2
3 best_para_svm_linear.fit(X_train_k,y_train_k)
4 SVM_value_y_linear = best_para_svm_linear.predict(X_test_k)
5 print('Training score: {:.3f}'.format(best_para_svm_linear.score(X_train_k, y_train_k)))
6 print('Testing score: {:.3f}'.format(best_para_svm_linear.score(X_test_k, y_test_k)))
7

```

[LibSVM]Training score: 0.819  
 Testing score: 0.760



In [129]:

```
1 print(classification_report(SVM_value_y_linear,y_test_k))
```

	precision	recall	f1-score	support
0.0	1.00	0.76	0.86	250
1.0	0.00	0.00	0.00	0
accuracy			0.76	250
macro avg	0.50	0.38	0.43	250
weighted avg	1.00	0.76	0.86	250


In [130]:

```
1 print(confusion_matrix(y_pred = SVM_value_y_linear,y_true = y_test_k))
```

```
[[190  0]
 [ 60  0]]
```

In [131]:

```
1 %matplotlib inline
2
3 heatmap = mglearn.tools.heatmap(
4     confusion_matrix(y_pred = SVM_value_y_linear, y_true = y_test_k), xlabel = 'Predict
5     yticklabels=['No Default','Default'], cmap = "RdYlGn", fmt = "%d")
6 plt.title("Confusion matrix (poly)")
7 plt.gca().invert_yaxis()
```



In [132]:

```
1 Kernel_linear_precision_score=precision_score(y_test_k,best_para_svm_linear.predict(X_t
2 print('Precision score : {:.2f} '.format(precision_score(y_test_k,best_para_svm_linear.
3
```

Precision score : 0.00

In [133]:

```
1 Kernel_linear_f1_score=f1_score(y_test_k, best_para_svm_linear.predict(X_test_k))
2 print('f1 Score : {:.2f} '.format(Kernel_linear_f1_score))
3
```

f1 Score : 0.00

In [134]:

```
1 Kernel_linear_recall_score=recall_score(y_test_k, best_para_svm_linear.predict(X_test_k))
2 print('Recall score : {:.2f} '.format(Kernel_linear_recall_score))
3
```

Recall score : 0.00

In [135]:

```
1 Summary_Kernelized_linear= {'Type': 'Kernalized linear ', 'Train Score': best_para_svm_linear.score(X_train_k, y_train_k)*100,
2                             'Testing Score':best_para_svm_linear.score(X_test_k, y_test_k)*100,
3                             'f1 Score':f1_score(y_test_k, best_para_svm_linear.predict(X_test_k))};
```

In [136]:

```
1 Summary_Kernelized_linear
```

Out[136]:

```
{'Type': 'Kernalized linear ',
 'Train Score': 81.86666666666666,
 'Testing Score': 76.0,
 'f1 Score': 0.0}
```

```
1 <b>5.Decision Tree Classification.</b>
```

In [137]:

```
1 from sklearn.tree import DecisionTreeClassifier
2
3 dtree_training = []
4 dtree_testing = []
5
6 for depth in range(1,20):
7     dtree = DecisionTreeClassifier(max_depth = depth, random_state = 0)
8     dtree.fit(X_train, y_train)
9     dtree_training.append(dtree.score(X_train, y_train))
10    dtree_testing.append(dtree.score(X_test, y_test))
```

In [138]:


```

1 xvalues = range(1,20)
2
3 plt.plot(xvalues, dtree_training, label = 'Training Score', c = 'g')
4 plt.plot(xvalues, dtree_testing, label = 'Testing Score', c = 'r')
5 plt.xlabel('Max_depth')
6 plt.ylabel('Accuracy')
7 plt.legend()

```

Out[138]:

&lt;matplotlib.legend.Legend at 0x2331167e580&gt;



Based on above graph max\_depth = 7

In [139]:

```

1 dtree = DecisionTreeClassifier(max_depth = 7)
2
3 dtree.fit(X_train, y_train)
4 dtree_value_y = dtree.predict(X_test)
5
6 print('Training score: {:.3f}'.format(dtree.score(X_train, y_train)))
7 print('Testing score: {:.3f}'.format(dtree.score(X_test, y_test)))

```

Training score: 0.834

Testing score: 0.826


## Decision Tree Feature Selection

In [140]:

```
1 feat_importance = pd.Series(dtree.feature_importances_, index=X.columns)
2 feat_importance.nlargest(X.shape[1]).plot(kind='barh', figsize=(10, 10))
```

Out[140]:

<AxesSubplot:>



**PAY\_O: Repayment status in September, 2005 is the most important feature.**

**Cross validation scores for Decision Tree Classifier**

In [141]:

```
1 dtree_CV_scores = cross_val_score(dtree, X_train,y_train, cv = 7)
2
3 pd.DataFrame({"Cross-Validation scores": dtree_CV_scores})
```

Out[141]:

Cross-Validation scores	
0	0.814619
1	0.820840
2	0.822029
3	0.813317
4	0.817984
5	0.813939
6	0.821717

In [142]:

```
1 print("cross-validation score is : {:.2f}".format(dtree_CV_scores.mean()))
```

cross-validation score is : 0.82

## Grid Search with Decision Tree Classifier

In [143]:

```
1 dtree.get_params()
```

Out[143]:

```
{'ccp_alpha': 0.0,
 'class_weight': None,
 'criterion': 'gini',
 'max_depth': 7,
 'max_features': None,
 'max_leaf_nodes': None,
 'min_impurity_decrease': 0.0,
 'min_impurity_split': None,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'presort': 'deprecated',
 'random_state': None,
 'splitter': 'best'}
```

In [144]:

```

1 param_grid_dtree = {'max_depth': range(1,20), 'criterion':['gini','entropy'], 'min_sample
2 CV_dtrees = GridSearchCV(estimator = dtree, cv = 7, param_grid = param_grid_dtree , ver
3 GS_results_dtrees = CV_dtrees.fit(X_train, y_train)
4 best_parameters_dtrees = CV_dtrees.best_params_
5 print(best_parameters_dtrees)
6
7

```

Fitting 7 folds for each of 1824 candidates, totalling 12768 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 0.5s
[Parallel(n_jobs=-1)]: Done 328 tasks    | elapsed: 3.6s
[Parallel(n_jobs=-1)]: Done 828 tasks    | elapsed: 11.8s
[Parallel(n_jobs=-1)]: Done 1528 tasks   | elapsed: 30.5s
[Parallel(n_jobs=-1)]: Done 2428 tasks   | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 3528 tasks   | elapsed: 1.9min
[Parallel(n_jobs=-1)]: Done 4634 tasks   | elapsed: 2.9min
[Parallel(n_jobs=-1)]: Done 5384 tasks   | elapsed: 3.7min
[Parallel(n_jobs=-1)]: Done 6234 tasks   | elapsed: 4.6min
[Parallel(n_jobs=-1)]: Done 7944 tasks   | elapsed: 5.4min
[Parallel(n_jobs=-1)]: Done 9594 tasks   | elapsed: 6.9min
[Parallel(n_jobs=-1)]: Done 10744 tasks  | elapsed: 8.4min
[Parallel(n_jobs=-1)]: Done 11994 tasks  | elapsed: 10.2min
[Parallel(n_jobs=-1)]: Done 12768 out of 12768 | elapsed: 11.4min finished

```

```
{'criterion': 'entropy', 'max_depth': 4, 'min_samples_leaf': 48}
```

In [145]:

```

1 print("Best score : Decision Tree grid search ")
2 round(GS_results_dtrees.best_score_,2)

```

Best score : Decision Tree grid search

Out[145]:

0.82

In [146]:

```

1 print("Best parameters : Decision Tree grid search ")
2 GS_results_dtrees.best_params_

```

Best parameters : Decision Tree grid search

Out[146]:

```
{'criterion': 'entropy', 'max_depth': 4, 'min_samples_leaf': 48}
```

In [147]:

```

1 best_dtree = DecisionTreeClassifier(criterion='entropy', max_depth=4, min_samples_leaf=
2
3 best_dtree.fit(X_train, y_train)
4 dtree_value_y = best_dtree.predict(X_test)
5
6 print('Training score: {:.3f}'.format(best_dtree.score(X_train, y_train)))
7 print('Testing score: {:.3f}'.format(best_dtree.score(X_test, y_test)))

```

Training score: 0.825

Testing score: 0.829

In [148]:

```

1 print(classification_report(y_pred =dtree_value_y, y_true = y_test))

```

	precision	recall	f1-score	support
0.0	0.85	0.95	0.90	5950
1.0	0.66	0.35	0.46	1550
accuracy			0.83	7500
macro avg	0.76	0.65	0.68	7500
weighted avg	0.81	0.83	0.81	7500

In [149]:

```

1 print(confusion_matrix(y_pred = dtree_value_y, y_true = y_test))

```


```

[[5674  276]
 [1003  547]]

```

In [150]:

```
1 %matplotlib inline
2
3 heatmap = mglearn.tools.heatmap(
4     confusion_matrix(y_pred = dtree_value_y, y_true = y_test), xlabel = 'Predicted label',
5     ylabel='True label', xticklabels = ['No Default','Default'],yticklabels=['No Default','Default'])
6 plt.title("Confusion matrix (Decision Tree)'No Default','Default'")
7
8 plt.gca().invert_yaxis()
```



In [151]:

```
1 dtree_precision_score=precision_score(y_test, best_dtree.predict(X_test))
2 print('Precision score : {:.2f} '.format(dtree_precision_score))
```

Precision score : 0.66

In [152]:

```
1 dtree_recall_score = recall_score(y_test, best_dtree.predict(X_test))
2 print('Recall score : {:.2f} '.format(dtree_recall_score))
```

Recall score : 0.35

In [153]:

```
1 dtree_f1_score = f1_score(y_test,best_dtree.predict(X_test))
2 print('f1 score : {:.2f} '.format(dtree_f1_score))
```

f1 score : 0.46



In [154]:

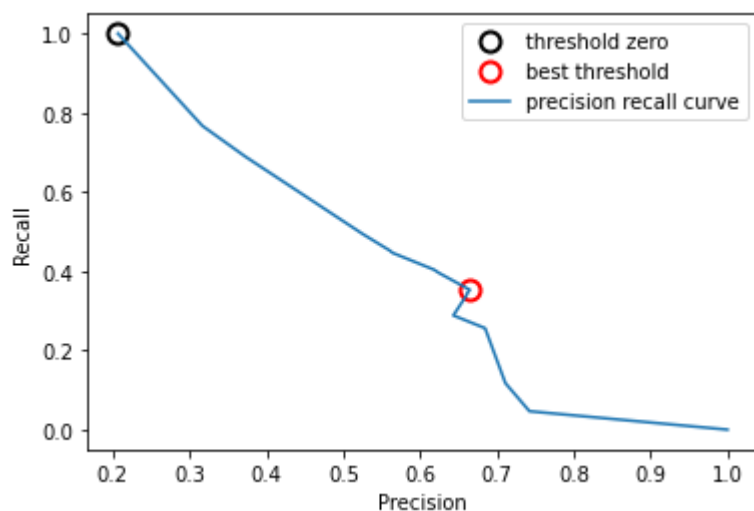
```

1 import mglearn
2 %matplotlib inline
3 %matplotlib inline
4
5 precision, recall, thresholds = precision_recall_curve(y_test, best_dtree.predict_proba
6
7 plt.plot(precision[close_zero], recall[close_zero], 'o', markersize=10,
8          label="threshold zero", fillstyle="none", c='k', mew=2)
9 plt.plot(dtree_precision_score, dtree_recall_score, 'o', markersize=10,
10          label="best threshold", fillstyle="none", c='r', mew=2)
11
12 plt.plot(precision, recall, label="precision recall curve")
13 plt.xlabel("Precision")
14 plt.ylabel("Recall")
15 plt.legend(loc="best")

```

Out[154]:

&lt;matplotlib.legend.Legend at 0x233113e2b80&gt;



In [155]:

```

1 Summary_dtree= {'Type': 'Decision Tree', 'Train Score': best_dtree.score(X_train, y_train),
2                 'Testing Score':best_dtree.score(X_test, y_test)*100,
3                 'f1 Score':f1_score(y_test, best_dtree.predict(X_test))};

```

In [156]:

```
1 Summary_dtree
```

Out[156]:

```
{'Type': 'Decision Tree',  
 'Train Score': 82.45333333333333,  
 'Testing Score': 82.94666666666667,  
 'f1 Score': 0.46101980615254945}
```

## Comparing All Models

In [157]:

```
1 Summary_Knn
```

Out[157]:

```
{'Type': 'K-nearest Neighbors (KNN) Classification Model',  
 'Training Score': 100.0,  
 'Testing Score': 81.0,  
 'f1 Score': 0.35256701499318494}
```

In [158]:

```
1 Summary_Logistic
```

Out[158]:

```
{'Type': 'Logistic Regression',  
 'Train Score': 81.05333333333333,  
 'Testing Score': 81.65333333333334,  
 'f1 Score': 0.2993890020366599}
```

In [159]:

```
1 Summary_lin_SVM
```

Out[159]:

```
{'Type': 'Linear SVM',  
 'Train Score': 79.63555555555556,  
 'Testing Score': 80.28,  
 'f1 Score': 0.12329579134558386}
```

In [160]:

```
1 Summary_Kernelized_rbf
```

Out[160]:

```
{'Type': 'Kernalized rbf ',  
 'Train Score': 81.86666666666666,  
 'Testing Score': 76.0,  
 'f1 Score': 0.0}
```

In [161]:

```
1 Summary_Kernelized_linear
```

Out[161]:

```
{'Type': 'Kernalized linear ',  
'Train Score': 81.86666666666666,  
'Testing Score': 76.0,  
'f1 Score': 0.0}
```

In [162]:

```
1 Summary_Kernelized_poly
```

Out[162]:

```
{'Type': 'Kernalized poly ',  
'Train Score': 85.6,  
'Testing Score': 76.8,  
'f1 Score': 0.17142857142857143}
```

In [163]:

```
1 Summary_dtree
```

Out[163]:

```
{'Type': 'Decision Tree',  
'Train Score': 82.45333333333333,  
'Testing Score': 82.94666666666667,  
'f1 Score': 0.46101980615254945}
```

In [164]:

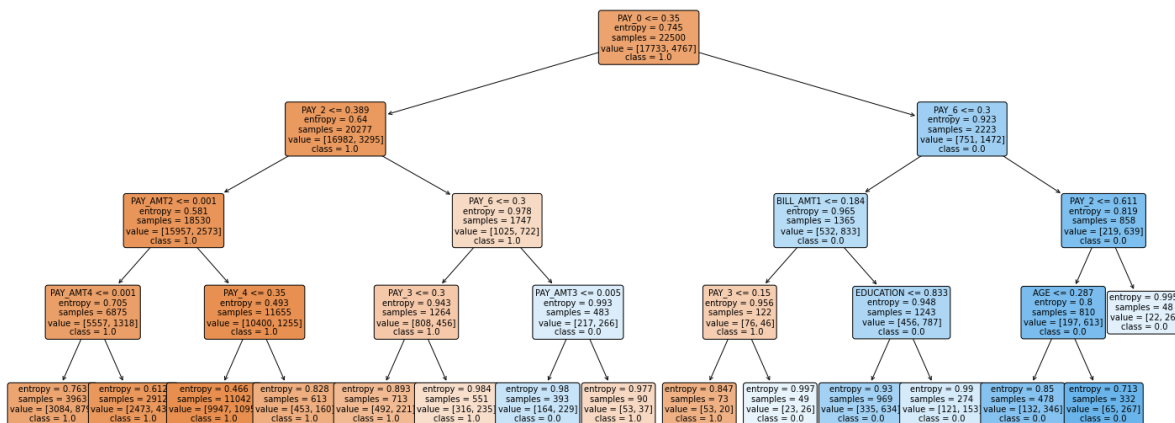
```
1 X = dcc.drop('Default_Payment',axis =1)  
2 y = dcc['Default_Payment']  
3 X_train_org, X_test_org, y_train, y_test = train_test_split(X, y, random_state = 0)  
4  
5
```

In [165]:

```

1 from sklearn.tree import DecisionTreeClassifier, plot_tree
2 dtree_feature_names = X.columns
3 dtree_target_names = y.unique().astype(str).tolist()
4 plt.figure(figsize=(25,10))
5 plot_tree(best_dtree,
6           feature_names = dtree_feature_names,
7           class_names = dtree_target_names,
8           filled = True,
9           rounded = True, fontsize=10)
10
11 plt.savefig('tree_visualization.png')
12
13 plt.show()
14

```



1) From above summary we can state that Decision Tree is the best classification model as the training and test scores are 82.45% and 82.94% respectively with a f1 score of 0.46 which is highest compared to other candidates, since we have different model predicting the same thing, f1 score is a good choice for this purpose as it depends heavily on how imbalanced our training dataset is.

2) Since the target variable in our dataset is highly skewed towards one particular category, thus decision tree performs fairly better than other classification models in this case.

3) PAY\_O: Repayment status in September, 2005 is the most important feature

## End of Project 1 : Classification

*Initials:*

-rp